

Using Hadoop: Best Practices

Casey Stella

March 14, 2012

Table of Contents

Introduction

Background

Using Hadoop Professionally

Indexing

Performance

Staying Sane

Testing

Debugging

State of Big Data and Hadoop

Conclusion

Introduction

- ▶ Hi, I'm Casey
 - ▶ I work at Explorys
 - ▶ I work with Hadoop and the Hadoop ecosystem daily

Introduction

- ▶ Hi, I'm Casey
 - ▶ I work at Explorys
 - ▶ I work with Hadoop and the Hadoop ecosystem daily
- ▶ I'm going to talk about some of the best practices that I've seen
 - ▶ Some of these are common knowledge
 - ▶ Some of these don't show up until you've been up 'til 3AM debugging a problem.

Introduction

- ▶ Hi, I'm Casey
 - ▶ I work at Explorys
 - ▶ I work with Hadoop and the Hadoop ecosystem daily
- ▶ I'm going to talk about some of the best practices that I've seen
 - ▶ Some of these are common knowledge
 - ▶ Some of these don't show up until you've been up 'til 3AM debugging a problem.
- ▶ These are my opinions and not necessarily the opinions of my employer.

The Lay of the Land – The Bad

- ▶ There are two APIs, prefer the mapred package
 - ▶ The mapreduce and the mapred packages
 - ▶ mapred is deprecated, but still preferred
 - ▶ Hortonworks just kind of screwed up

The Lay of the Land – The Bad

- ▶ There are two APIs, prefer the mapred package
 - ▶ The mapreduce and the mapred packages
 - ▶ mapred is deprecated, but still preferred
 - ▶ Hortonworks just kind of screwed up
- ▶ The Pipes interface is really poorly implemented and very slow
- ▶ HDFS currently has a single point of failure

The Lay of the Land – The Good

- ▶ Hortonworks is actively working on Map-Reduce v2
 - ▶ This means other distributed computing models
 - ▶ Included in 0.23

The Lay of the Land – The Good

- ▶ Hortonworks is actively working on Map-Reduce v2
 - ▶ This means other distributed computing models
 - ▶ Included in 0.23
- ▶ HDFS is dramatically faster in 0.23
 - ▶ Socket communication is made more efficient
 - ▶ Smarter checksumming

Indexing

- ▶ Hadoop is a batch processing system, but you need realtime access

Indexing

- ▶ Hadoop is a batch processing system, but you need realtime access
- ▶ Options are
 - ▶ Roll your own (Jimmy Lin talks about how one might serve up inverted indices in Chapter 3)

Indexing

- ▶ Hadoop is a batch processing system, but you need realtime access
- ▶ Options are
 - ▶ Roll your own (Jimmy Lin talks about how one might serve up inverted indices in Chapter 3)
 - ▶ Use an open source indexing infrastructure, like Katta

Indexing

- ▶ Hadoop is a batch processing system, but you need realtime access
- ▶ Options are
 - ▶ Roll your own (Jimmy Lin talks about how one might serve up inverted indices in Chapter 3)
 - ▶ Use an open source indexing infrastructure, like Katta
 - ▶ Serve them directly from HDFS with an on-disk index aka Hadoop MapFiles

Indexing

- ▶ Hadoop is a batch processing system, but you need realtime access
- ▶ Options are
 - ▶ Roll your own (Jimmy Lin talks about how one might serve up inverted indices in Chapter 3)
 - ▶ Use an open source indexing infrastructure, like Katta
 - ▶ Serve them directly from HDFS with an on-disk index aka Hadoop MapFiles
 - ▶ Serve them through HBase or Cassandra

Indexing

- ▶ Hadoop is a batch processing system, but you need realtime access
- ▶ Options are
 - ▶ Roll your own (Jimmy Lin talks about how one might serve up inverted indices in Chapter 3)
 - ▶ Use an open source indexing infrastructure, like Katta
 - ▶ Serve them directly from HDFS with an on-disk index aka Hadoop MapFiles
 - ▶ Serve them through HBase or Cassandra
 - ▶ If data permits, push them to a database

Indexing

- ▶ Hadoop is a batch processing system, but you need realtime access
- ▶ Options are
 - ▶ Roll your own (Jimmy Lin talks about how one might serve up inverted indices in Chapter 3)
 - ▶ Use an open source indexing infrastructure, like Katta
 - ▶ Serve them directly from HDFS with an on-disk index aka Hadoop MapFiles
 - ▶ Serve them through HBase or Cassandra
 - ▶ If data permits, push them to a database
- ▶ Katta can serve up both Lucene indices and Mapfiles

Indexing

- ▶ Hadoop is a batch processing system, but you need realtime access
- ▶ Options are
 - ▶ Roll your own (Jimmy Lin talks about how one might serve up inverted indices in Chapter 3)
 - ▶ Use an open source indexing infrastructure, like Katta
 - ▶ Serve them directly from HDFS with an on-disk index aka Hadoop MapFiles
 - ▶ Serve them through HBase or Cassandra
 - ▶ If data permits, push them to a database
- ▶ Katta can serve up both Lucene indices and Mapfiles
- ▶ Indexing is hard, be careful.

Performance Considerations

- ▶ Setup and teardown costs, so keep the HDFS block size large
- ▶ Mappers, Reducers and Combiners have memory constraints
- ▶ Transmission costs dearly
 - ▶ Use Snappy, LZO, or (soon) LZ4 compression at every phase
 - ▶ Serialize your objects tightly (e.g. not using Java Serialization)
 - ▶ Key/values emitted from the map phase had better be linear with a **small** constant..preferably below 1

Performance Considerations

► Strategies

- Intelligent use of the combiners
- Use Local Aggregation in the mapper to emit a more complex value. (you already know this)
- Ensure that all components of your keys are necessary in the sorting logic. If any are not, push them into the value.

¹http://hadoop.apache.org/common/docs/current/mapred_tutorial.html#Profiling

²<http://hadoop.apache.org/common/docs/current/voidya.html> ►

Performance Considerations

- ▶ Strategies
 - ▶ Intelligent use of the combiners
 - ▶ Use Local Aggregation in the mapper to emit a more complex value. (you already know this)
 - ▶ Ensure that all components of your keys are necessary in the sorting logic. If any are not, push them into the value.
- ▶ Profile **JobConf.setProfileEnabled(boolean)**¹
- ▶ Use Hadoop Vaidya²

¹http://hadoop.apache.org/common/docs/current/mapred_tutorial.html#Profiling

²<http://hadoop.apache.org/common/docs/current/vaidya.html>

Unit/Integration Testing Methodologies

- ▶ First off, do it.
- ▶ Unit test individual mappers, reducers, combiners and partitioners
 - ▶ Actual unit tests. This will help debugging, I promise.
 - ▶ Design components so that dependencies can be injected via polymorphism when testing

Unit/Integration Testing Methodologies

- ▶ First off, do it.
- ▶ Unit test individual mappers, reducers, combiners and partitioners
 - ▶ Actual unit tests. This will help debugging, I promise.
 - ▶ Design components so that dependencies can be injected via polymorphism when testing
- ▶ Minimally verify that keys
 - ▶ Can be serialized and deserialized
 - ▶ *hashCode()* is sensible (Remember: the *hashCode()* for enum is not stable across different JVMs instances)
 - ▶ *compareTo()* is reflexive, symmetric and jives with *equals()*
- ▶ Integration test via single user mode hadoop

Quality Assurance Testing

- ▶ The output of processing large amounts of data is often large
- ▶ Verify statistical properties

Quality Assurance Testing

- ▶ The output of processing large amounts of data is often large
- ▶ Verify statistical properties
 - ▶ If statistical tests fit within Map Reduce, then use MR
 - ▶ If not, then sample the dataset with MR and verify with R, Python or whatever.
- ▶ Do outlier analysis and thresholding based QA

Debugging Methodologies

- ▶ Better to catch it at the unit test level
- ▶ If you can't, I suggest the following technique
 - ▶ Investigatory map reduce job to find the data causing the issue.
 - ▶ Single point if you're lucky, if not then a random sample using reservoir sampling
 - ▶ Take the data and integrate it into a unit test.

Debugging Methodologies

- ▶ Better to catch it at the unit test level
- ▶ If you can't, I suggest the following technique
 - ▶ Investigatory map reduce job to find the data causing the issue.
 - ▶ Single point if you're lucky, if not then a random sample using reservoir sampling
 - ▶ Take the data and integrate it into a unit test.
- ▶ **DO NOT**
 - ▶ Use print statements to debug unless you're sure of the scope.
 - ▶ Use counters where the group or name count grows more than a fixed amount.

Debugging Methodologies

- ▶ Better to catch it at the unit test level
- ▶ If you can't, I suggest the following technique
 - ▶ Investigatory map reduce job to find the data causing the issue.
 - ▶ Single point if you're lucky, if not then a random sample using reservoir sampling
 - ▶ Take the data and integrate it into a unit test.
- ▶ **DO NOT**
 - ▶ Use print statements to debug unless you're sure of the scope.
 - ▶ Use counters where the group or name count grows more than a fixed amount.
- ▶ **DO**
 - ▶ Use a single counter in the actual job if the job doesn't finish
 - ▶ Use a map reduce job that outputs suspect input data into HDFS

Hadoop Opinions

- ▶ “We’re about a year behind Google” – Doug Cutting, Hadoop World

Hadoop Opinions

- ▶ “We’re about a year behind Google” – Doug Cutting, Hadoop World
- ▶ Giraph and Mahout are just not there yet

Hadoop Opinions

- ▶ “We’re about a year behind Google” – Doug Cutting, Hadoop World
- ▶ Giraph and Mahout are just not there yet
- ▶ HBase is getting there (Facebook is dragging HBase into being serious)

Hadoop Opinions

- ▶ “We’re about a year behind Google” – Doug Cutting, Hadoop World
- ▶ Giraph and Mahout are just not there yet
- ▶ HBase is getting there (Facebook is dragging HBase into being serious)
- ▶ Zookeeper is the real deal

Hadoop Opinions

- ▶ “We’re about a year behind Google” – Doug Cutting, Hadoop World
- ▶ Giraph and Mahout are just not there yet
- ▶ HBase is getting there (Facebook is dragging HBase into being serious)
- ▶ Zookeeper is the real deal
- ▶ Cassandra is cool, but eventual consistency is too hard to seriously consider.

Big Data

- ▶ We kind of went overboard w.r.t. Map Reduce
 - ▶ Easier than MPI, but really not as flexible.
 - ▶ Bringing distributed computing to the masses...meh, maybe the masses don't need it.
 - ▶ M.R. v2 opens up a broader horizon

Big Data

- ▶ We kind of went overboard w.r.t. Map Reduce
 - ▶ Easier than MPI, but really not as flexible.
 - ▶ Bringing distributed computing to the masses...meh, maybe the masses don't need it.
 - ▶ M.R. v2 opens up a broader horizon
- ▶ Data analysis is hard and often requires specialized skills
 - ▶ Enter a new breed: the data scientist
 - ▶ Stats + Computer Science + Domain knowledge
 - ▶ Often not a software engineer

Conclusion

- ▶ Thanks for your attention
- ▶ Follow me on twitter @casey_stella
- ▶ Find me at
 - ▶ <http://caseystella.com>
 - ▶ <https://github.com/cestella>
- ▶ P.S. If you dig this stuff, come work with me.